



Minimax Algorithm

Tic-Tac-Toe



Can we use the search algorithms seen so far to decide our next move in a tic-tac-toe match?



Tic-Tac-Toe

Tic-Tac-Toe



Can we use the search algorithms seen so far to decide our next move in a tic-tac-toe match?



We miss the idea that our adversaries are actively trying to defeat us!



Tic-Tac-Toe

Tic-Tac-Toe




Can we use the search algorithms seen so far to decide our next move in a tic-tac-toe match?



We miss the idea that our adversaries are actively trying to defeat us!



“To defeat your enemy, you must become your enemy.”



Tic-Tac-Toe

Competitive Environments


- Environments where two or more agents have conflicting goals
- Adversarial search problems



Tic-Tac-Toe

Two-Player Zero-Sum Games

- Deterministic, two-player, turn-taking, perfect information, zero-sum games
 - “perfect information” → “fully observable”
 - “zero-sum” → what is good for one player is just as bad for the other
- We will concentrate on such games because:
 - the state of a game is easy to represent
 - agents are usually restricted to a small number of actions whose effects are defined by precise rules



Two-Player Zero-Sum Games

- 1952: Arthur Samuel's checkers playing program



Figure 1. Playing checkers on the 701

1952

Two-Player Zero-Sum Games

- 1994: Chinook ended 40-year-reign of human champion Marion Tinsley using complete 8-piece endgame



Figure 2. Chinook team (August 1992). From left to right: Duane Szafron, Joe Culberson, Paul Lu, Brent Knight, Jonathan Schaeffer, Rob Lake, and Steve Sutphen. Our checkers expert, Norman Treloar, is missing.

1952

1994

Two-Player Zero-Sum Games

- 1997: IBM's Deep Blue beats world chess champion



Figure 3. Garry Kasparov in a 1997 game against Deep Blue.

1997

1952


1994

Two-Player Zero-Sum Games

- 2007: Chinook – the unbeatable checkers player



Figure 4. Chinook team (August 2012). From left to right Steve Sutphen, Duane Szafron, Jonathan Schaeffer, Rob Lake, and Paul Lu.




Two-Player Zero-Sum Games

- 2017: Google DeepMind's AlphaGo defeats Go champions



Figure 5. World champion Ke Jie struggles against AlphaGo, a product of Alphabet's (formerly Google's) subsidiary DeepMind.



Adversarial Search

- For you to play optimally, your opponent must play optimally



Adversarial Search

- Maximize the utility of your moves, minimize the utility of your opponent's moves
 - Utility function gives a numeric value for a terminal state
 - Example: win, loss, or draw, with values 1, -1, or 0

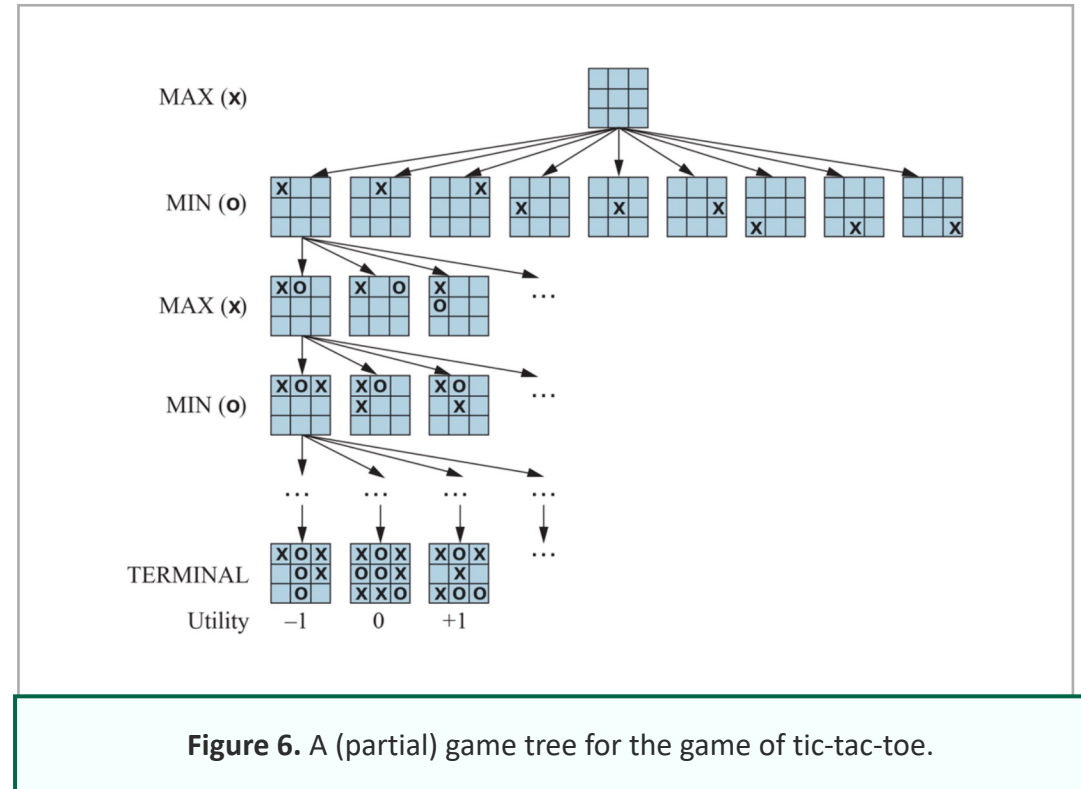



Figure 6. A (partial) game tree for the game of tic-tac-toe.

The Minimax Algorithm

- $\text{minimax}(\mathbf{s}) =$
 - $\text{utility}(\mathbf{s})$ for the MAX player if \mathbf{s} is terminal
 - $\max_{a \in \text{Actions}(\mathbf{s})} \text{minimax}(\text{move}(\mathbf{s}, a))$ if MAX is playing
 - $\min_{a \in \text{Actions}(\mathbf{s})} \text{minimax}(\text{move}(\mathbf{s}, a))$ if MIN is playing



The Minimax Algorithm - Pseudocode

```
(action) minimax(s):
```

```
    a, val = max_player(s)
```

```
    return a
```

```
(action,utility) max_player(s):
```

```
    if s is terminal:
```

```
        return null, utility_function(s)
```

```
    max_a, max_val = null,  $-\infty$ 
```

```
    for a in actions:
```

```
        b, val = min_player(move(s,a))
```

```
        if val > max_val:
```

```
            max_a, max_val = a, val
```

```
    return max_a, max_val
```

```
(action,utility) min_player(s):
```

```
    if s is terminal:
```

```
        return null, utility_function(s)
```

```
    min_a, min_val = null,  $\infty$ 
```

```
    for a in actions:
```

```
        b, val = max_player(move(s,a))
```

```
        if val < min_val:
```

```
            min_a, min_val = a, val
```

```
    return min_a, min_val
```

```
(utility) utility_function(s):
```

```
    return utility of state s for the max_player
```

```
(state) move(s,a):
```

```
    return state reached from s after taking action a
```

Knowledge Check 1



What will be the values of the nodes A, B, C, D, E, F, and G in the minimax search tree on the left?




Figure 8. Minimax search tree.

A

4, 8, 4, 6, 8, 1, 4

B

100, 100, 20, 10, 100, 2, 20

C

10, 10, 2, 10, 100, 2, 20

D

6, 6, 1, 6, 8, 1, 4

Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?


Alpha-Beta Pruning

- The number of game states is exponential in the depth of the tree
- No algorithm can completely eliminate the exponent, but we can sometimes cut it in half by pruning large parts of the tree that make no difference to the outcome
- The particular technique we will examine is called alpha–beta pruning


Alpha-Beta Pruning

- Alpha–beta pruning gets its name from the two extra parameters that describe bounds along the search path:
 - α = the value of the best choice we have found so far along the path for MAX (α = “at least”)
 - β = the value of the best choice we have found so far along the path for MIN (β = “at most”)


Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning




Alpha-Beta Pruning



Alpha-Beta Pruning



Alpha-Beta Pruning



Alpha-Beta Pruning


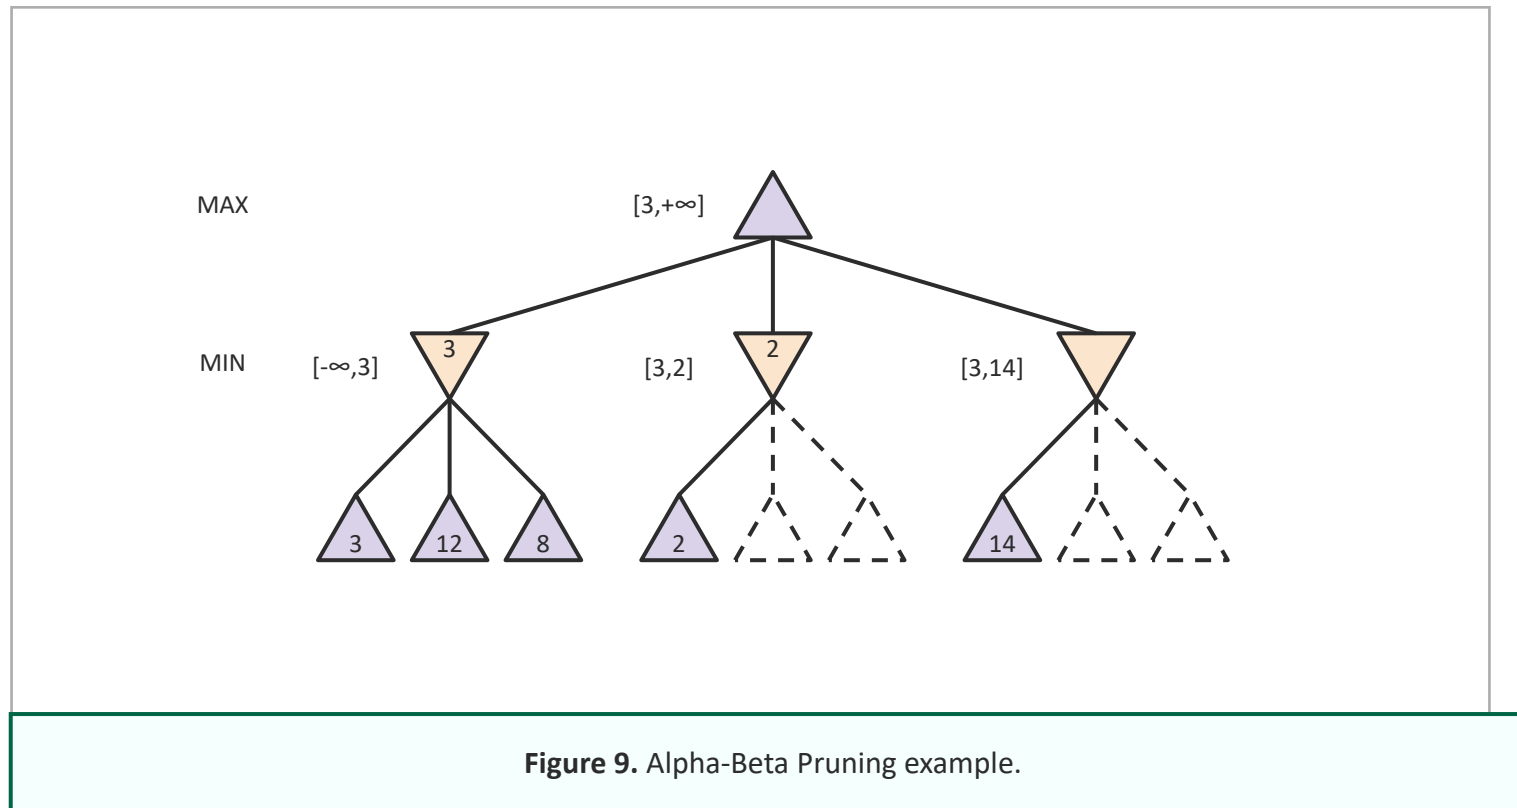


Figure 9. Alpha-Beta Pruning example.

Alpha-Beta Pruning



Alpha-Beta Pruning




Figure 9. Alpha-Beta Pruning example.

Alpha-Beta Pruning




Figure 9. Alpha-Beta Pruning example.

Alpha-Beta Pruning - Pseudocode

```
(action) alpha_beta(s):  
    a, val = max_player(s,  $-\infty$ ,  $+\infty$ )  
    return a  
  
(action, utility) max_player(s,  $\alpha$ ,  $\beta$ ):  
    if s is terminal:  
        return null, utility_function(s)  
    max_a, max_val = null,  $-\infty$   
    for a in actions:  
        b, val =  
min_player(move(s, a),  $\alpha$ ,  $\beta$ )  
        if val > max_val:  
            max_a, max_val = a, val  
             $\alpha$  = max( $\alpha$ , val)  
        if max_val  $\geq$   $\beta$ :  
            break  
    return max_a, max_val
```

```
(action, utility) min_player(s,  $\alpha$ ,  $\beta$ ):  
    if s is terminal:  
        return null, utility_function(s)  
    min_a, min_val = null,  $\infty$   
    for a in actions:  
        b, val = max_player(move(s, a),  $\alpha$ ,  $\beta$ )  
        if val < min_val:  
            min_a, min_val = a, val  
             $\beta$  = min( $\beta$ , val)  
        if min_val  $\leq$   $\alpha$ :  
            break  
    return min_a, min_val
```

Knowledge Check 2



Which of the internal nodes A-G and terminal nodes (a)-(h) will not be expanded when using Alpha-Beta pruning?




Figure 10. Minimax search tree.

A

E, (c), (d), (f), (h)

B

E, G, (c), (d), (f), (g), (h)

C

G, (d), (g), (h)

D

(d), (h)

Alpha-Beta Pruning Properties

- This pruning has no effect on minimax value computed for the root!
- Values of intermediate nodes might be wrong
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - Full search of complex games, e.g. chess, is still hopeless...



You have reached the end
of the lecture.



Image/Figure References

Figure 1. Playing checkers on the 701. Source: <https://www.ibm.com/ibm/history/ibm100/us/en/icons/ibm700series/impacts/>

Figure 2. Chinook team (August 1992). From left to right: Duane Szafron, Joe Culberson, Paul Lu, Brent Knight, Jonathan Schaeffer, Rob Lake, and Steve Sutphen. Our checkers expert, Norman Treloar, is missing. Source: <http://jonathanschaeffer.blogspot.com/2012/08/chinook-twenty-years-later.html>

Figure 3. Garry Kasparov in a 1997 game against Deep Blue. Source: <https://www.businessinsider.com/how-ibm-watson-is-transforming-healthcare-2015-7>

Figure 4. Chinook team (August 2012). From left to right Steve Sutphen, Duane Szafron, Jonathan Schaeffer, Rob Lake, and Paul Lu. Source:

<http://jonathanschaeffer.blogspot.com/2012/08/chinook-twenty-years-later.html>

Figure 5. World champion Ke Jie struggles against AlphaGo, a product of Alphabet's (formerly Google's) subsidiary DeepMind. Source: <https://www.theneweconomy.com/technology/alphago-defeats-world-go-champion-ke-jie>

Figure 6. A (partial) game tree for the game of tic-tac-toe. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.

Figure 7. Minimax search tree. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.

Figure 8. Minimax search tree. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.

Figure 9. Alpha-Beta Pruning example. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.

Figure 10. Minimax search tree. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.

Other images were purchased from Getty Images and with permission to use.