



Monte Carlo Tree Search

Alpha-Beta Pruning Weakness

- When the branching factor is high, minimax with Alpha-Beta pruning must run in "depth-limited" mode
 - Stop when maximum depth is reached, then use an heuristic function to evaluate the current state of the game
 - Example:
 - The branching factor of Go is 361
 - The search cannot go deeper than 4-5 moves
 - It is difficult to define a good evaluation function for Go
 - material value is not a strong indicator
 - most positions are in flux until the endgame

Monte Carlo Algorithms

- Randomized algorithms named after the Casino de Monte-Carlo in Monaco
- Example
 - Estimate the area of a circle with radius 1 by randomly sampling points within a 2×2 square and measuring their distance to the center of the square. Points whose distance is smaller than or equal to 1 are inside the circle.
 - $\text{Area} = 4 \times |\text{number of points in circle}| / |\text{number of sampled points}|$




Figure 1. Monte Carlo algorithms.

Knowledge Check 1



Monte Carlo simulations can be used for many applications. For instance, we can estimate the probability of having two people with the same birthday in a group of N people using the code on the right. For which value of N this probability is higher than 50%?

```
import random

def rand():
    return random.randint(1, 365)

def montecarlo(N, simulations=10000):
    count = 0
    for i in range(simulations):
        birthdays = set([rand() for _ in range(N)])
        if len(birthdays) < N:
            count += 1
    return count/simulations
```

A

13

B

23

C

33

D

43

Monte Carlo Tree Search

- The value of a state is estimated as the average utility over a number of simulations of complete games starting from the state
- A simulation (or playout) chooses moves first for one player, then for the other, repeating until a terminal position is reached
- At the terminal position, the rules of the game determine who has won or lost, and by what score
- For games in which the only outcomes are a win or a loss, “average utility” is the same as “win percentage”

Monte Carlo Tree Search

- Just choosing moves randomly during playouts is not good enough for most games
 - Use a playout policy to bias the moves towards good ones
 - Example
 - Prioritize capture moves in chess
- Starting multiple playouts from a state to choose the next move is not good enough for most games



MCTS uses a search tree and a selection policy to balance exploration and **exploitation**.

Monte Carlo Tree Search

Repeat:

01

Selection

Starting at the root of the search tree, choose moves guided by the selection policy until a leaf is reached.

02

Expansion

Grow the search tree by generating a new child for the selected node.

03

Simulation

Perform a playout from the newly generated node, choosing moves for both players according to the playout policy.

04

Back-Propagation

Use the result of the simulation to update the winning statistics of all search tree nodes going up to the root.

Selection Policy

- The value of a state is estimated as the average utility over a number of simulations of complete games starting from the state
- Upper confidence bounds applied to trees (UCT)
 - The policy ranks each possible move based on an upper confidence bound formula called UCB1

$$UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log N(\text{PARENT}(n))}{N(n)}}$$

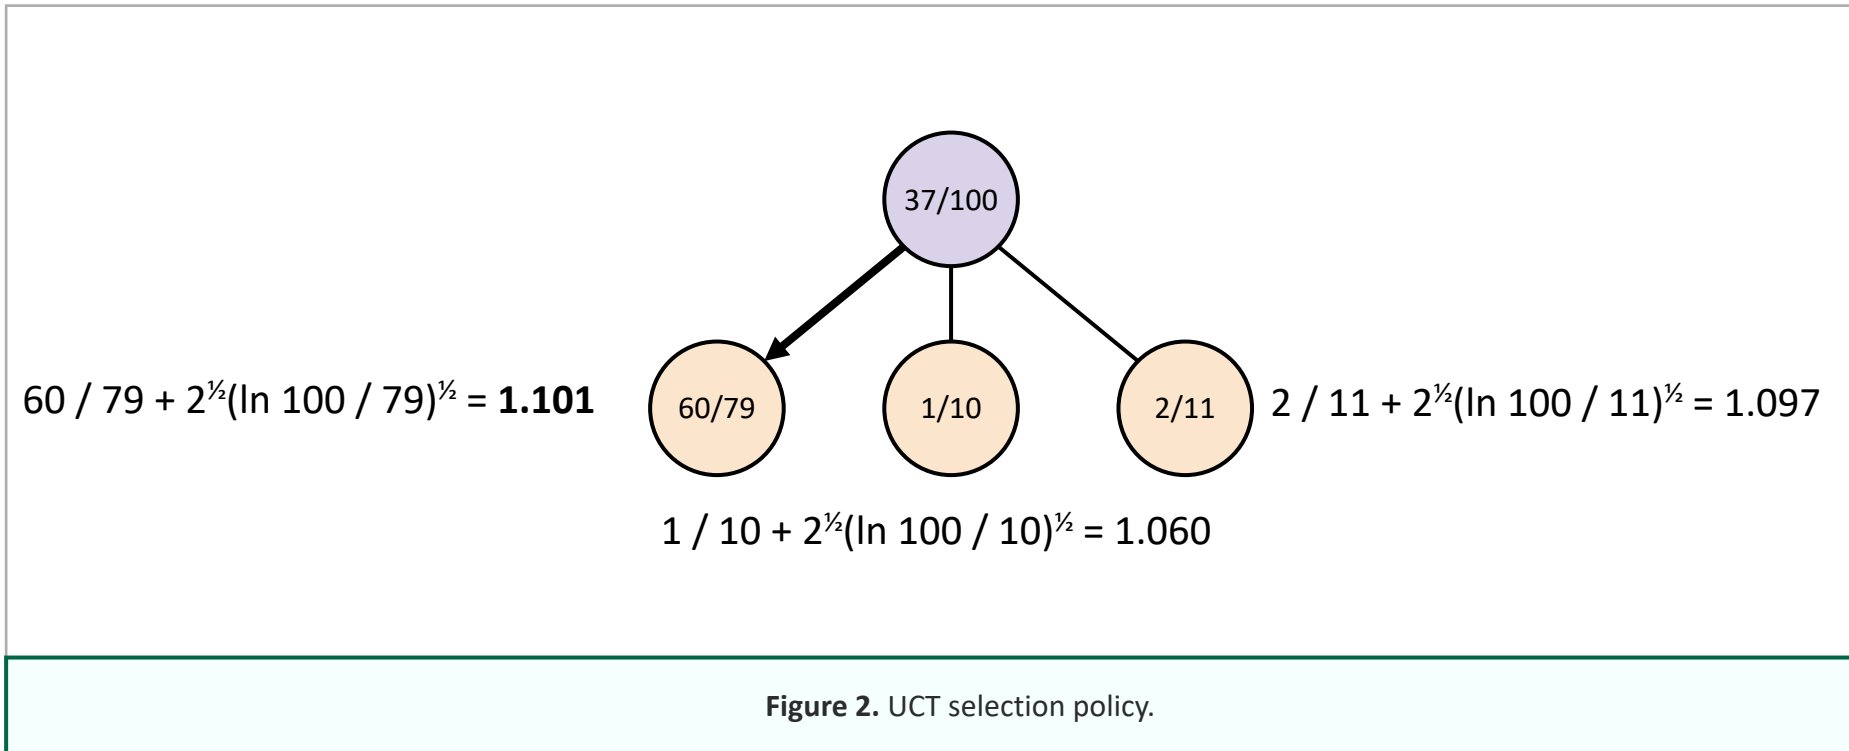
$U(n)$ is the total utility of all playouts through node n

$N(n)$ is the number of playouts through node n

p is the parent node of n

C is a constant that balances exploitation and exploration ($C \approx 2^{1/2}$)

Selection Policy



[illegible]

Monte Carlo Tree Search

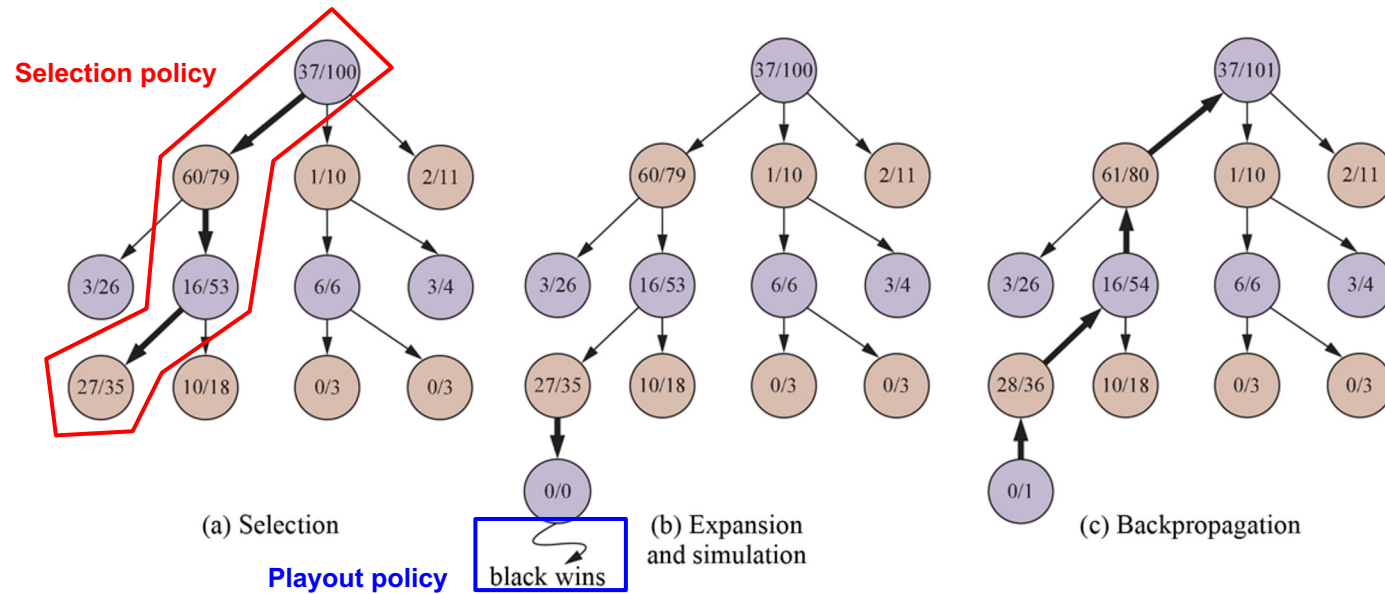


Figure 4. One iteration of the process of choosing a move with MCTS using the UCT selection policy.

Knowledge Check 2



When the average game length is high, the playout simulation might take too long to reach a terminal state. For those problems, one possible solution is stopping the simulation early and:

A

picking a random winner.

B

considering this playout a draw.

C

using a heuristic function to estimate the utility.

D

restarting the playout with a different random seed.



You have reached the end
of the lecture.



Image/Figure References

Figure 1. Monte Carlo algorithms.

Figure 2. UCT selection policy. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.

Figure 3. One iteration of the process of choosing a move with MCTS using the UCT selection policy. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.

Figure 4. One iteration of the process of choosing a move with MCTS using the UCT selection policy. Source: Russell & Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Pearson, 2021.