



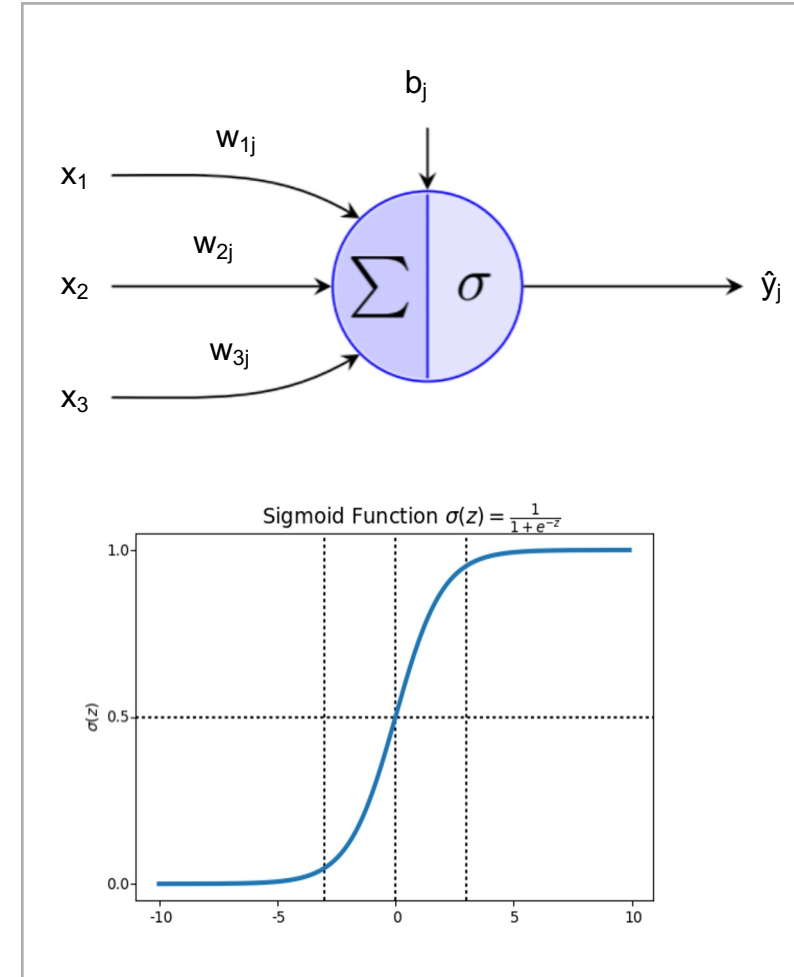
Multilayer Perceptron

Perceptron

- Input may have multiple values
- Output may have multiple values
 - Multiple parallel Perceptrons
 - Input is the same for all Perceptrons
- Activation function – sigmoid
 - $\sigma' = \sigma(1 - \sigma)$

$$\hat{\mathbf{y}} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$$

$$\mathcal{L} = (1/2) \sum_k (\hat{y}_k - y_k)^2$$



$$\frac{\partial \mathcal{L}}{\partial w_{ij}} = \frac{\partial}{\partial w_{ij}} \left(\frac{1}{2} \right) \sum (\hat{y}_k - y_k)^2 = \frac{\partial}{\partial w_{ij}} \left(\frac{1}{2} \right) (\hat{y}_j - y_j)^2 =$$

$$(\hat{y}_j - y_j) \frac{\partial}{\partial w_{ij}} (\hat{y}_j - y_j) = (\hat{y}_j - y_j) \frac{\partial}{\partial w_{ij}} \sigma(\mathbf{x} \mathbf{w}_{*j} + b_j) - y_j =$$

$$(\hat{y}_j - y_j) \sigma(\hat{y}_j) (1 - \sigma(\hat{y}_j)) \frac{\partial}{\partial w_{ij}} \mathbf{x} \mathbf{w}_{*j} = (\hat{y}_j - y_j) \sigma(\hat{y}_j) (1 - \sigma(\hat{y}_j)) x_i$$

$$\frac{\partial \mathcal{L}}{\partial b_j} = (\hat{y}_j - y_j) \sigma(\hat{y}_j) (1 - \sigma(\hat{y}_j))$$

Multilayer Perceptron (MLP)

- Consecutive layers of parallel Perceptrons
- Introduces the concept of hidden layer

$$\mathbf{h} = \sigma(W_1\mathbf{x} + \mathbf{b}_1)$$

$$\hat{\mathbf{y}} = \sigma(W_2\mathbf{h} + \mathbf{b}_2)$$

$$\mathcal{L} = (1/2) \sum_k (\hat{y}_k - y_k)^2$$

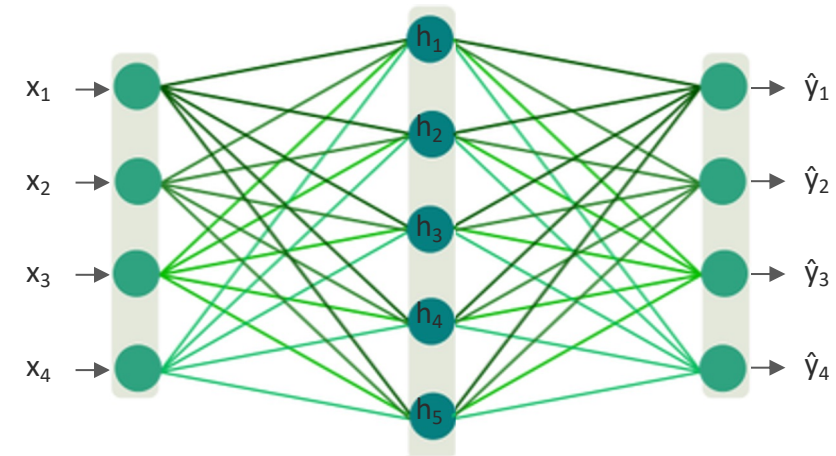
Number of parameters:

$$|W_1| = |\mathbf{x}| * |\mathbf{h}|$$

$$|\mathbf{b}_1| = |\mathbf{h}|$$

$$|W_2| = |\mathbf{h}| * |\mathbf{y}|$$

$$|\mathbf{b}_2| = |\mathbf{y}|$$



[1] Rosenblatt, F. (1961). Principles of Neurodynamics: Perceptrons and the Theory of Brain Mechanisms. Spartan Books, Washington DC

[2] Rumelhart, D., Hinton, G. & Williams, R. (1986). Learning representations by back-propagating errors. Nature 323, 533-536

Knowledge Check 1



What happens if we do not use an activation function in between layers of parallel Perceptrons?

A

Nothing changes

B

The network loses representation capacity, so we will need to add many more layers to achieve similar results

C

No matter how many layers we use, the network will have the same capacity of a single-layer network

Why do we need activation functions?

- Last layer: regularize output
- Intermediate layers: handle nonlinearities
 - Consecutive linear layer are equivalent to a single layer

$$\mathbf{h} = \mathbf{W}_1\mathbf{x} + \mathbf{b}_1$$

$$\hat{\mathbf{y}} = \mathbf{W}_2\mathbf{h} + \mathbf{b}_2$$

$$\hat{\mathbf{y}} = \mathbf{W}_2(\mathbf{W}_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2$$

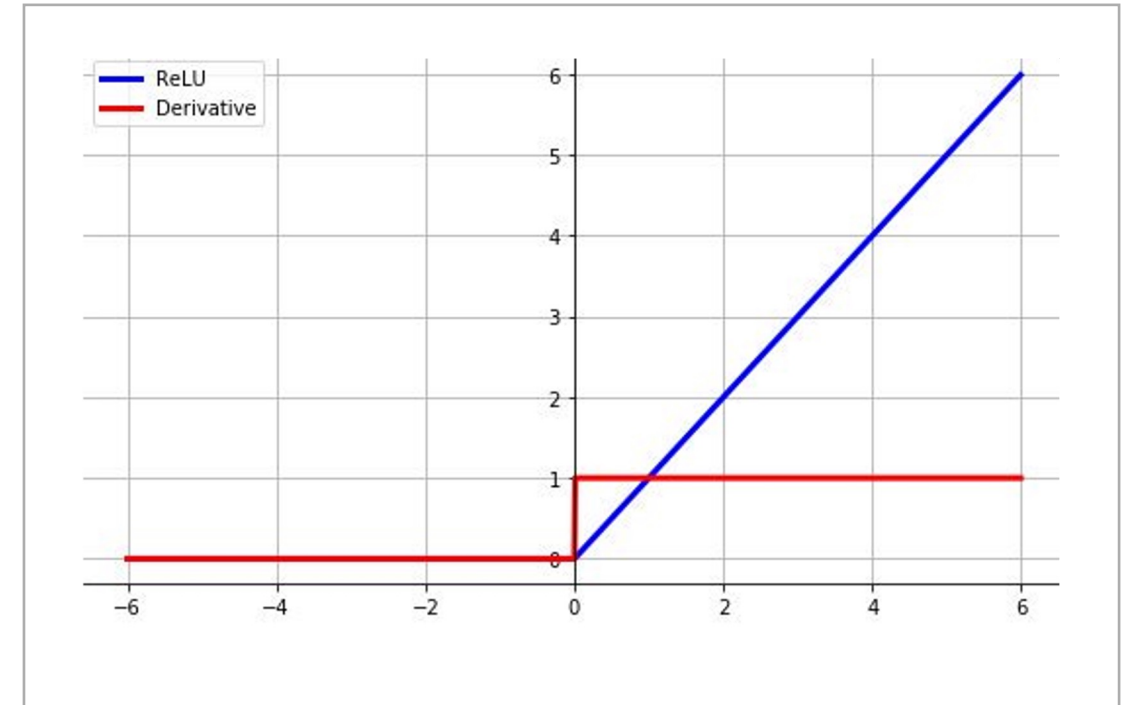
$$\hat{\mathbf{y}} = \mathbf{W}_2\mathbf{W}_1\mathbf{x} + \mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2$$

$$\mathbf{W} = \mathbf{W}_2\mathbf{W}_1$$

$$\mathbf{b} = \mathbf{W}_2\mathbf{b}_1 + \mathbf{b}_2$$

Activation Functions

- Rectified Linear Unit (ReLU)
- $\text{ReLU}(z) = \max(0, z)$
 - $\text{ReLU}'(z) = 1$ if $z \geq 0$
 - $\text{ReLU}'(z) = 0$ if $z < 0$
- Simplicity
- Linear behavior
- Avoids saturation



[1] Hahnloser, R. & Seung, H. (2001). Permitted and Forbidden Sets in Symmetric Threshold-Linear Networks. Neural Information Processing Systems

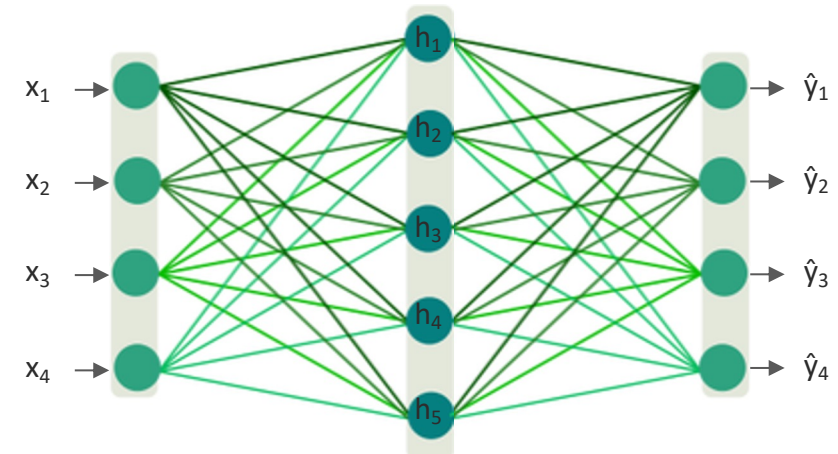
[2] Glorot, X., Bordes, A. & Bengio, Y. (2011). Deep Sparse Rectifier Neural Networks. International Conference on Artificial Intelligence and Statistics

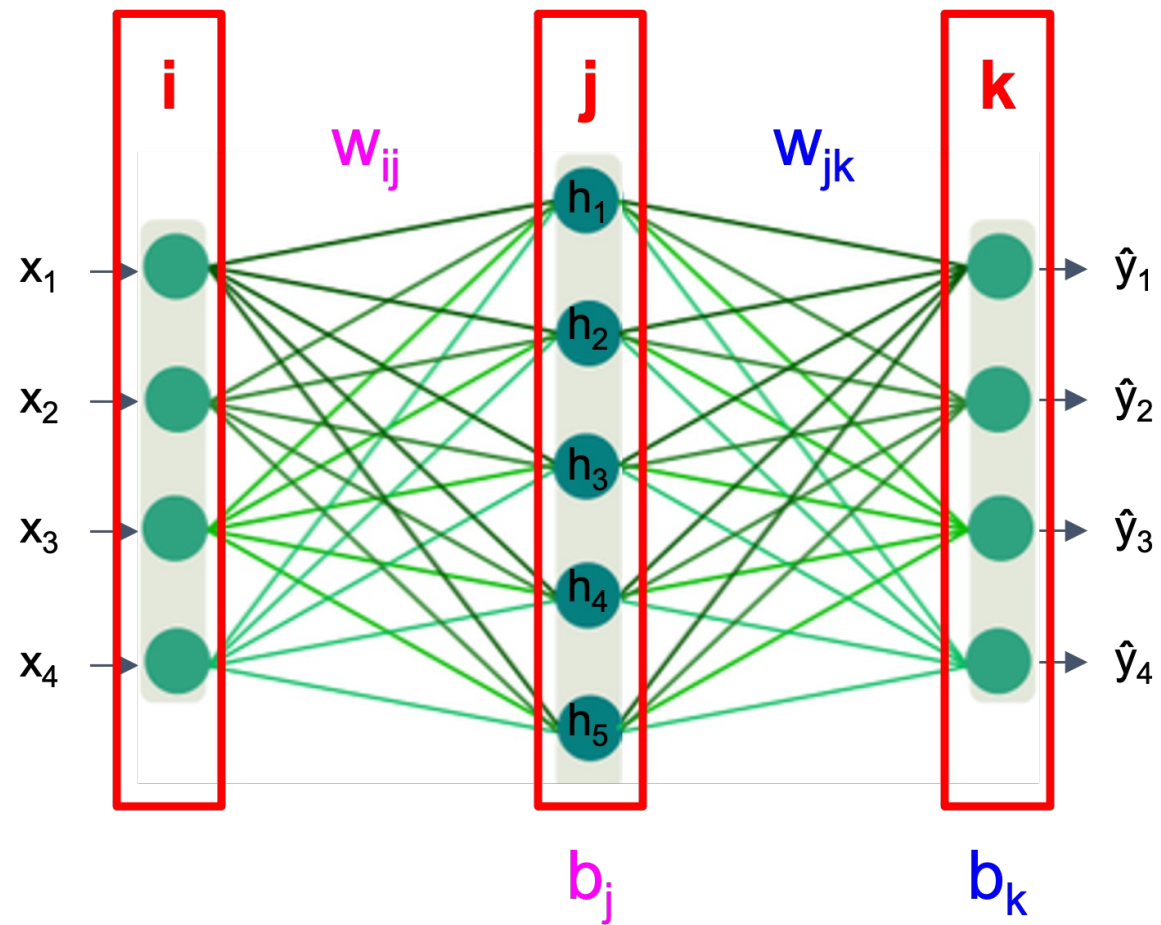
Updated Multilayer Perceptron

$$\mathbf{h} = \text{ReLU}(W_1 \mathbf{x} + \mathbf{b}_1)$$

$$\hat{\mathbf{y}} = W_2 \mathbf{h} + \mathbf{b}_2$$

$$\mathcal{L} = (1/2) \sum_k (\hat{y}_k - y_k)^2$$





$$\mathbf{w}_{jk}^{t+1} = \mathbf{w}_{jk}^t - \lambda \nabla_{\mathbf{w}[jk]} = \mathbf{w}_{jk}^t - \lambda \frac{\partial \mathcal{L}}{\partial \mathbf{w}_{jk}}$$

$$\mathbf{b}_k^{t+1} = \mathbf{b}_k^t - \lambda \nabla_{\mathbf{b}[k]} = \mathbf{b}_k^t - \lambda \frac{\partial \mathcal{L}}{\partial \mathbf{b}_k}$$

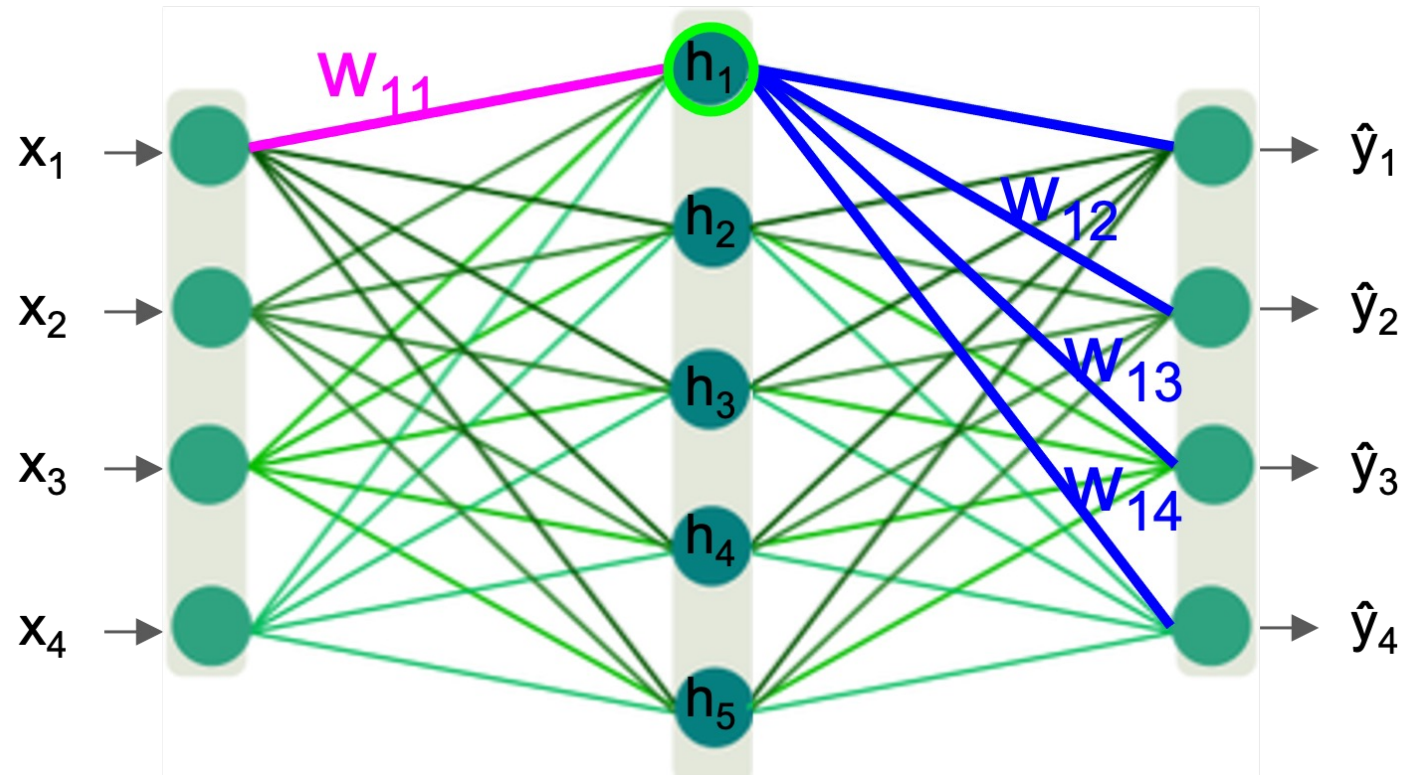
$$\delta_k = (\hat{y}_k - y_k)$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}_{jk}} = (\hat{y}_k - y_k) h_j = \delta_k h_j$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{b}_k} = (\hat{y}_k - y_k) = \delta_k$$

$$w_{ij}^{t+1} = w_{ij}^t - \lambda \nabla_{w[ij]} = w_{ij}^t - \lambda \frac{\partial \mathcal{L}}{\partial w_{ij}}$$

$$b_j^{t+1} = b_j^t - \lambda \nabla_{b[j]} = b_j^t - \lambda \frac{\partial \mathcal{L}}{\partial b_j}$$



$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}_{ij}} = \frac{\partial}{\partial \mathbf{w}_{ij}} (1/2) \sum_k (\hat{y}_k - y_k)^2 = (1/2) \sum_k \frac{\partial}{\partial \mathbf{w}_{ij}} (\hat{y}_k - y_k)^2 =$$

$$\sum_k (\hat{y}_k - y_k) \frac{\partial}{\partial \mathbf{w}_{ij}} \hat{y}_k = \sum_k \delta_k \frac{\partial}{\partial \mathbf{w}_{ij}} \mathbf{h} \mathbf{w}_{*k} + \mathbf{b}_k = \sum_k \delta_k \frac{\partial}{\partial \mathbf{w}_{ij}} \mathbf{h}_j \mathbf{w}_{jk} =$$

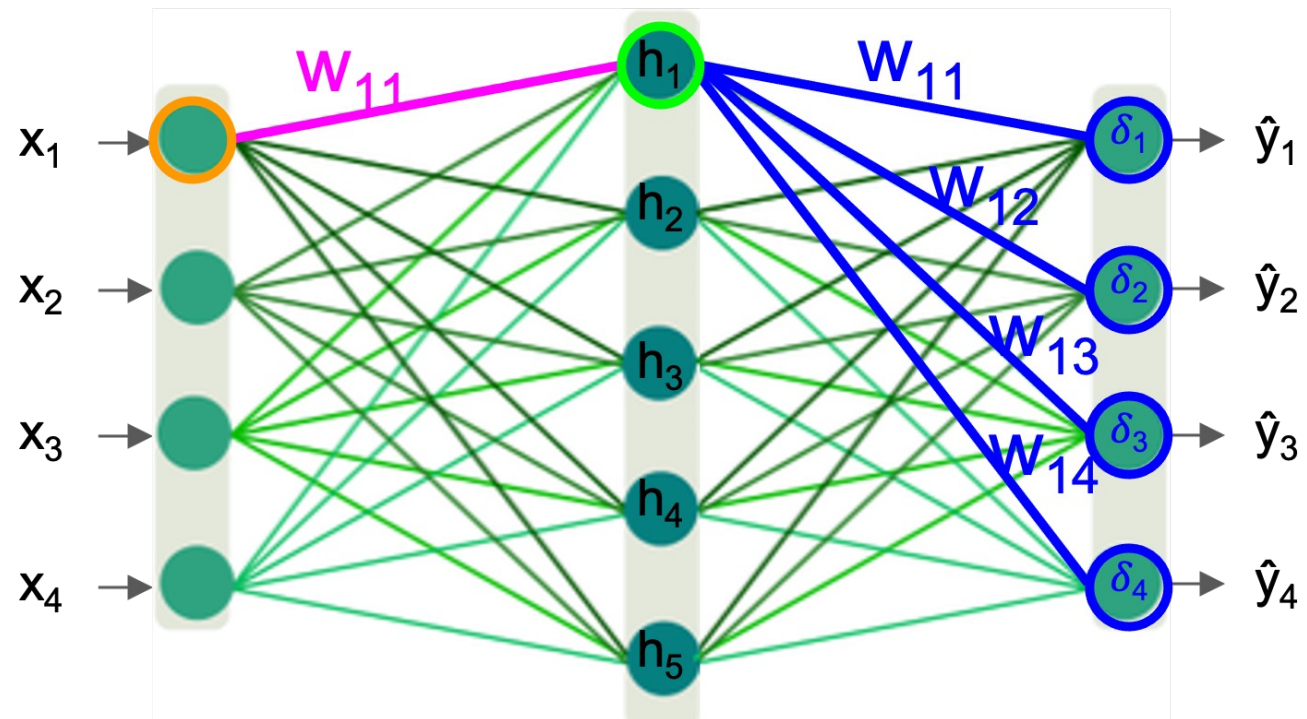
$$\sum_k \delta_k \frac{\partial}{\partial h_j} \mathbf{h}_j \mathbf{w}_{jk} \frac{\partial}{\partial \mathbf{w}_{ij}} \mathbf{h}_j = \sum_k \delta_k \mathbf{w}_{jk} \frac{\partial}{\partial \mathbf{w}_{ij}} \text{ReLU}(\mathbf{x} \mathbf{w}_{*j} + \mathbf{b}_j) =$$

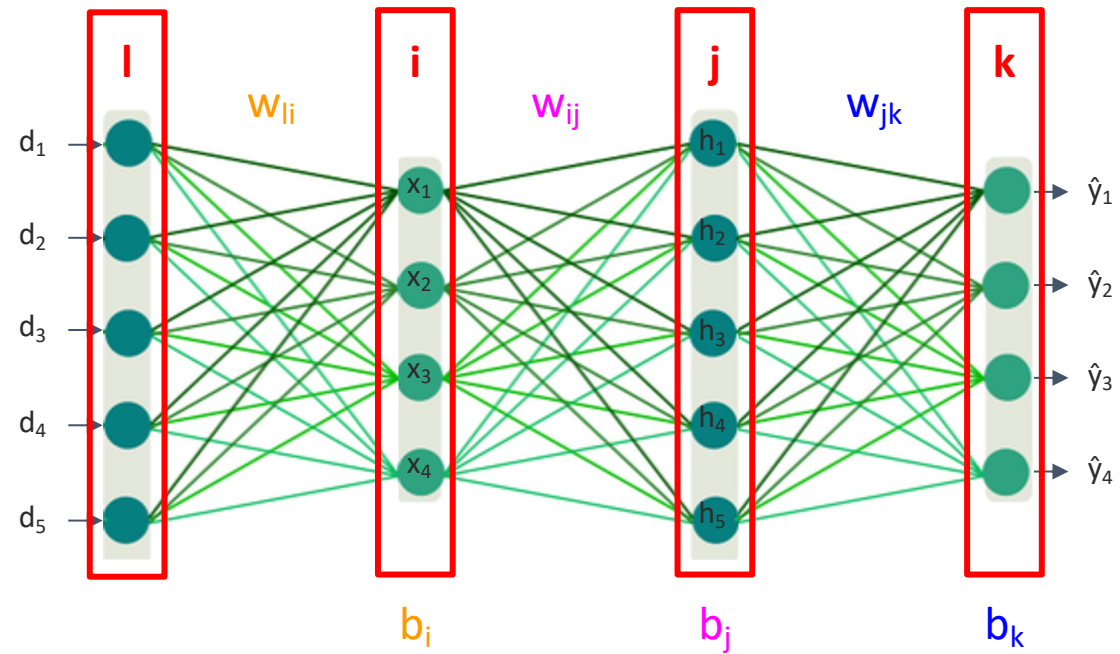
$$\sum_k \delta_k \mathbf{w}_{jk} \text{ReLU}'(h_j) \frac{\partial}{\partial \mathbf{w}_{ij}} \mathbf{x} \mathbf{w}_{*j} + \mathbf{b}_j = \text{ReLU}'(h_j) \mathbf{x}_i \sum_k \delta_k \mathbf{w}_{jk}$$

$$\delta_j = \begin{cases} \sum_k \delta_k w_{jk} & \text{if } h_j \geq 0 \\ 0 & \text{if } h_j < 0 \end{cases}$$

$$\frac{\partial \mathcal{L}}{\partial w_{ij}} = \delta_j x_i$$

$$\frac{\partial \mathcal{L}}{\partial b_j} = \delta_j$$





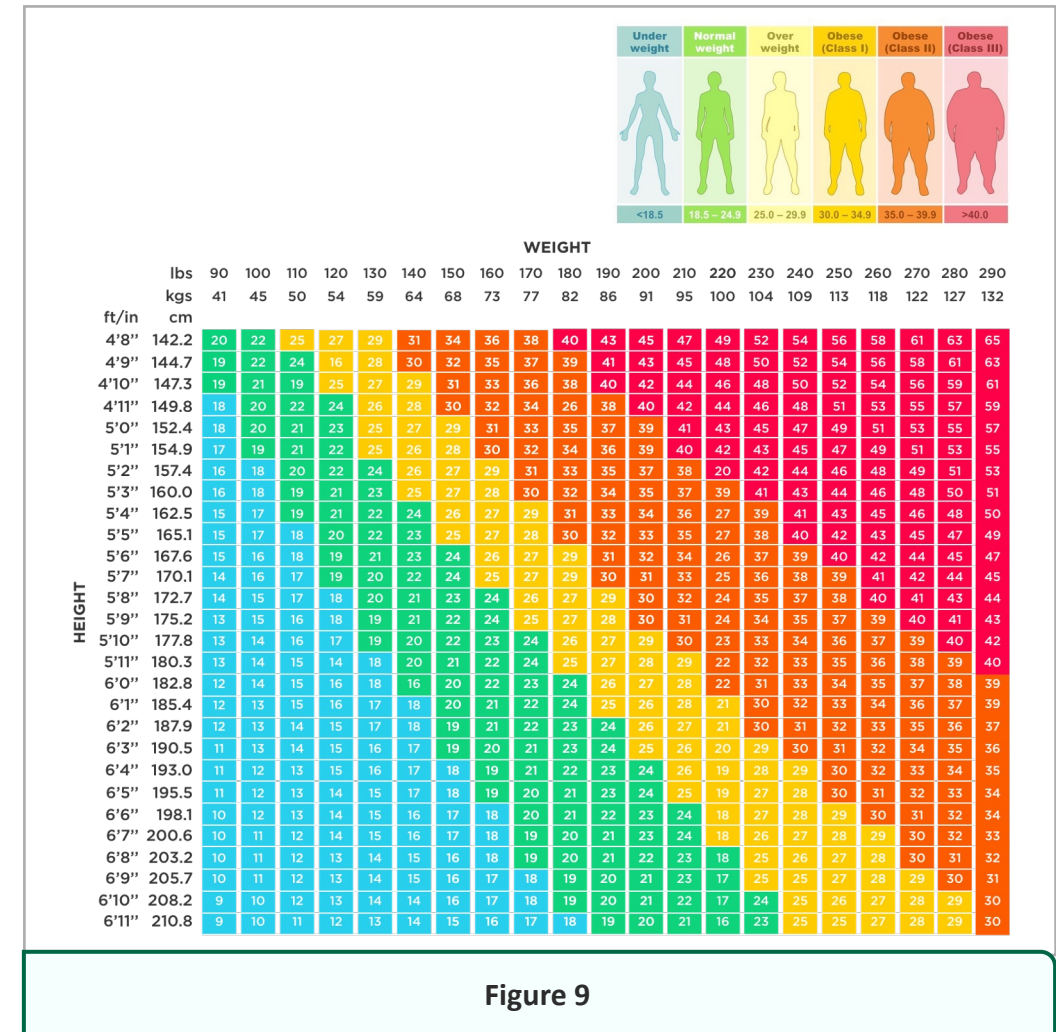
$$\delta_i = \begin{cases} \sum_j \delta_j w_{ij} & \text{if } x_i \geq 0 \\ 0 & \text{if } x_i < 0 \end{cases}$$

$$\frac{\partial \mathcal{L}}{\partial w_{li}} = \delta_i d_l$$

$$\frac{\partial \mathcal{L}}{\partial b_i} = \delta_i$$

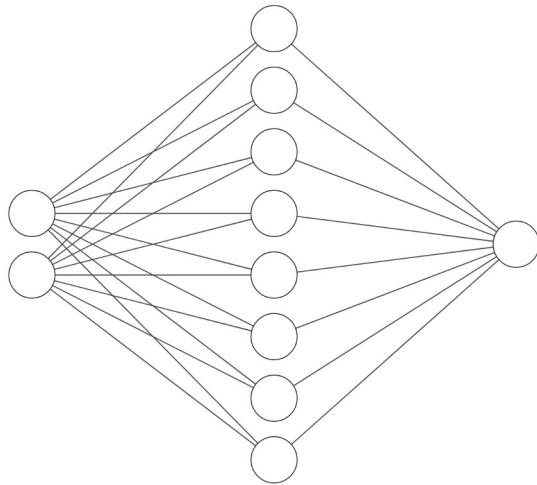
Non-Linear Regression – Multivariate Data

- $x = (\text{height, weight}) = (x_0, x_1)$
- $y = \text{BMI}$



Non-Linear Regression – Multivariate Data

- Let's use a 2-layer MLP with 8 hidden units and ReLU activation in the first layer
- $h_j = \text{ReLU}(w_{0j} * x_0 + w_{1j} * x_i + b_j)$
- $\hat{y} = w_0 h_0 + w_1 h_1 + \dots + w_7 h_7 + b$



- One gradient per weight:

$$\delta = (\hat{y} - y)$$

$$\frac{\partial \mathcal{L}}{\partial b} = \nabla_b = \delta$$

$$\frac{\partial \mathcal{L}}{\partial w_j} = \nabla_{w[j]} = \delta h_j$$

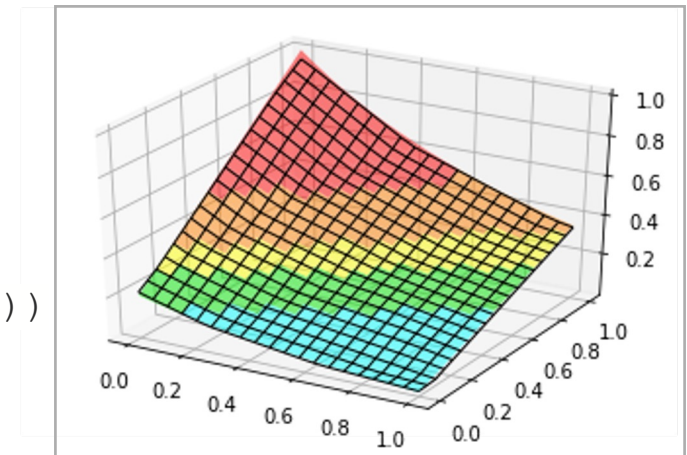
$$\delta_j = \begin{cases} \delta w_j & \text{if } h_j \geq 0 \\ 0 & \text{if } h_j < 0 \end{cases}$$

$$\frac{\partial \mathcal{L}}{\partial b_j} = \nabla_{b[j]} = \delta_j$$

$$\frac{\partial \mathcal{L}}{\partial w_{ij}} = \nabla_{w[i,j]} = \delta_j x_i$$

MLP - Gradient Descent

```
def gradient_descent_step(X, Y, W0, b0, W1, b1, learning_rate):  
    H_ = np.matmul(X, np.transpose(W0)) + b0  
    H = np.maximum(H_, 0)  
    Y_ = np.matmul(H, np.transpose(W1)) + b1  
  
    delta_Y = Y_ - Y  
    W1_grad = np.matmul(np.transpose(delta_Y), H) / H.shape[0]  
    b1_grad = np.mean(delta_Y, axis=0)  
  
    delta_H = np.multiply(np.matmul(delta_Y, W1), np.where(H_ >= 0.0, 1.0, 0.0))  
    W0_grad = np.matmul(np.transpose(delta_H), X) / X.shape[0]  
    b0_grad = np.mean(delta_H, axis=0)  
  
    W0_ = W0 - learning_rate * W0_grad  
    b0_ = b0 - learning_rate * b0_grad  
    W1_ = W1 - learning_rate * W1_grad  
    b1_ = b1 - learning_rate * b1_grad  
  
    return W0_, b0_, W1_, b1_
```



Common Practice #3: Weight Initialization

- Randomly initialize weights, otherwise learning is impossible
- Common approaches:
 - Normal distribution $\rightarrow \mathcal{N}(0, \sigma)$, where $\sigma = \sqrt{2 / (\text{\#in} + \text{\#out})}$
 - Uniform distribution $\rightarrow U(-a, a)$, where $a = \sqrt{6 / (\text{\#in} + \text{\#out})}$
 - \#in and \#out are the number of inputs and outputs in a weight tensor
- Transparent to the user in modern APIs

Knowledge Check 2



When using a 2-layer MLP for solving a problem, how many hidden units should we use?

A

As many as possible, so we can maximize the performance of the network

B

As little as possible, so we can minimize the computational cost

C

This number must be defined randomly, as there is no logical way to define it

D

This number must be at least as large as the input size



You have reached the end
of the lecture.

