



Autoencoders

What is an Autoencoder?

- A type of neural network used for learning compressed encodings of input data.
- There are 2 parts to an autoencoder: (1) Encoder and (2) Decoder
- Encoder (g): Takes the input data (x) and compresses it into a smaller form. Let's call this smaller form, $z = g(x)$.
- Decoder (f): Takes our compressed data, $z = g(x)$, and tries to decode it to as close as possible as the original input x .
- We'll call the re-generated input made by the decoder, $x' = f(z)$.

What is an Autoencoder?




Figure 1. Feed forward autoencoder.

Knowledge Check 1



Which of the following methods can be used for dimensionality reduction?

I) PCA II) LDA III) SVM IV) Autoencoder

A

I, II and III

B

I, III and IV

C

I, II and IV

D

II, III and IV

What is an Autoencoder?




Figure 2. Convolutional autoencoder.

What are the Uses of Autoencoders?

- Data Compression: For example, you can compress images into a smaller format.
- Dimensionality reduction - reducing the number of features from d to $k < d$. You can then use this smaller group of features for your supervised learning task, or you could use them for visualization if $k=1,2$, or 3 .
- Denoising: Removing “noise” from data.
- Unsupervised network initialization

Autoencoder Application: Feature Extraction

- You can use the “latent space representation” learned in the autoencoder for a different task such as classification.
- Step 1: Remove the decoder
- Step 2: Put your new layers for classification after the latent space, or “bottleneck” layer.
- Step 3: You can now perform your classification task.

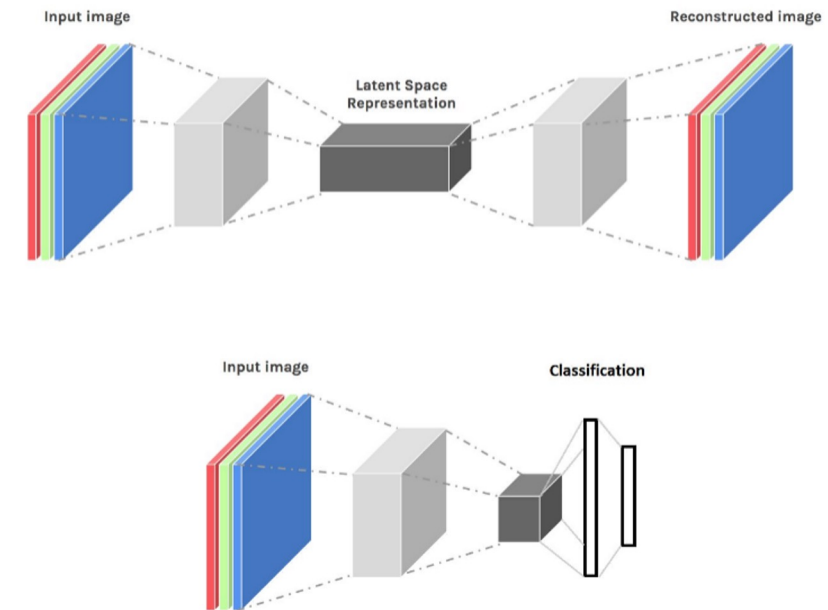


Figure 3. Autoencoder application: Feature extraction.

Face Autoencoder




Figure 4. Face autoencoder.

Upsampling in Keras

https://www.tensorflow.org/api_docs/python/tf/keras/layers/UpSampling2D

```
tf.keras.layers.UpSampling2D(  
    size=(2, 2),  
    data_format=None,  
    interpolation='nearest',  
    **kwargs  
)
```

Args	
size	Int, or tuple of 2 integers. The upsampling factors for rows and columns.
data_format	A string, one of channels_last (default) or channels_first. The ordering of the dimensions in the inputs. channels_last corresponds to inputs with shape (batch_size, height, width, channels) while channels_first corresponds to inputs with shape (batch_size, channels, height, width). It defaults to the image_data_format value found in your Keras config file at ~/.keras/keras.json. If you never set it, then it will be "channels_last".
interpolation	A string, one of "area", "bicubic", "bilinear", "gaussian", "lanczos3", "lanczos5", "mitchellcubic", "nearest".
Input shape	
4D tensor with shape:	
<ul style="list-style-type: none">• If data_format is "channels_last": (batch_size, rows, cols, channels)• If data_format is "channels_first": (batch_size, channels, rows, cols)	
Output shape	
4D tensor with shape:	
<ul style="list-style-type: none">• If data_format is "channels_last": (batch_size, upsampled_rows, upsampled_cols, channels)• If data_format is "channels_first": (batch_size, channels, upsampled_rows, upsampled_cols)	

Figure 5. Upsampling in Keras.

Transposed Convolutions in Keras

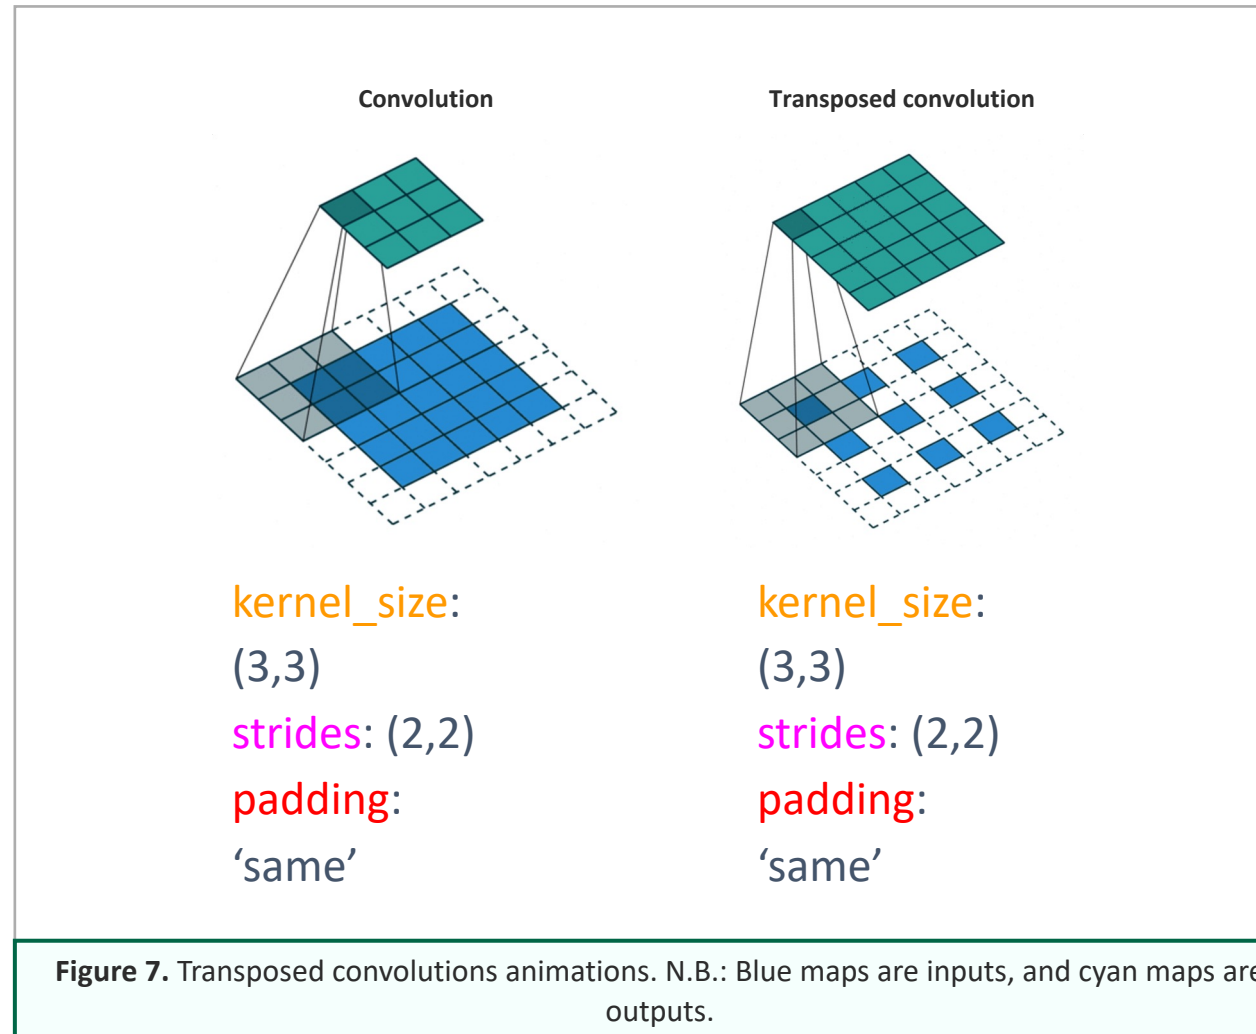
https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2DTranspose

```
tf.keras.layers.Conv2DTranspose(
    filters,
    kernel_size,
    strides=(1, 1),
    padding='valid',
    output_padding=None,
    data_format=None,
    dilation_rate=(1, 1),
    activation=None,
    use_bias=True,
    kernel_initializer='glorot_uniform',
    bias_initializer='zeros',
    kernel_regularizer=None,
    bias_regularizer=None,
    activity_regularizer=None,
    kernel_constraint=None,
    bias_constraint=None,
    **kwargs
)
```

Args	
filters	Integer, the dimensionality of the output space (i.e. the number of output filters in the convolution).
kernel_size	An integer or tuple/list of 2 integers, specifying the height and width of the 2D convolution window. Can be a single integer to specify the same value for all spatial dimensions.
strides	An integer or tuple/list of 2 integers, specifying the strides of the convolution along the height and width. Can be a single integer to specify the same value for all spatial dimensions. Specifying any stride value != 1 is incompatible with specifying any dilation_rate value != 1.
padding	one of "valid" or "same" (case-insensitive). "valid" means no padding. "same" results in padding with zeros evenly to the left/right or up/down of the input such that output has the same height/width dimension as the input.
Input shape	
4D tensor with shape: (batch_size, channels, rows, cols) if data_format='channels_first' or 4D tensor with shape: (batch_size, rows, cols, channels) if data_format='channels_last'.	
Output shape	
4D tensor with shape: (batch_size, filters, new_rows, new_cols) if data_format='channels_first' or 4D tensor with shape: (batch_size, new_rows, new_cols, filters) if data_format='channels_last'. rows and cols values might have changed due to padding. If output_padding is specified:	
<pre>new_rows = ((rows - 1) * strides[0] + kernel_size[0] - 2 * padding[0] + output_padding[0]) new_cols = ((cols - 1) * strides[1] + kernel_size[1] - 2 * padding[1] + output_padding[1])</pre>	
Returns	
A tensor of rank 4 representing <code>activation(conv2dtranspose(inputs, kernel) + bias)</code> .	

Figure 6. Transposed convolutions in Keras.

How Transposed Convolutions Work?





You have reached the end
of the lecture.



Image/Figure References

Figure 3. Autoencoder application: Feature extraction.

Figure 4. Face autoencoder.

Figure 5. Unsampling in Keras. Source: https://www.tensorflow.org/api_docs/python/tf/keras/layers/UpSampling2D

Figure 6. Transposed convolutions in Keras. Source: https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2DTranspose

Figure 7. Transposed convolutions animations. N.B.: Blue maps are inputs, and cyan maps are outputs. Source: https://github.com/vdumoulin/conv_arithmetic

Figure 8. Deepfakes with one-encoder & two-decoders. Source: Nguyen, T., Nguyen, C., Nguyen, D., Nguyen, D. & Nahavandi, S. (2020). Deep Learning for Deepfakes Creation and Detection: A Survey. arXiv abs/1909.11573.

Figure 9. Deepfakes with one-encoder & two-decoders. Source: Nguyen, T., Nguyen, C., Nguyen, D., Nguyen, D. & Nahavandi, S. (2020). Deep Learning for Deepfakes Creation and Detection: A Survey. arXiv abs/1909.11573.

Figure 10. Deepfakes with one-encoder & two-decoders. Source: Nguyen, T., Nguyen, C., Nguyen, D., Nguyen, D. & Nahavandi, S. (2020). Deep Learning for Deepfakes Creation and Detection: A Survey. arXiv abs/1909.11573.

Figure 11. Deepfakes with one-encoder & two-decoders. Source: Nguyen, T., Nguyen, C., Nguyen, D., Nguyen, D. & Nahavandi, S. (2020). Deep Learning for Deepfakes Creation and Detection: A Survey. arXiv abs/1909.11573.